

1 Automated Residual Plot Assessment with the R Package 2 autovi and the Shiny App autovi.web

3 Weihao Li^{1,2}, Dianne Cook¹, Emi Tanaka², Susan VanderPlas³ and Klaus
4 Ackermann¹

5 Monash University, The Australian National University and University of
6 Nebraska

Summary

7 Visual assessment of residual plots is a common approach for diagnosing linear
models, but it relies on manual evaluation, which does not scale well and can
lead to inconsistent decisions across analysts. The lineup protocol, which embeds
the observed plot among null plots, can reduce subjectivity but requires even
more human effort. In today's data-driven world, such tasks are well-suited for
automation. We present a new R package that uses a computer vision model to
automate the evaluation of residual plots. An accompanying Shiny app is provided
for ease of use. Given a sample of residuals, the model predicts a visual signal
strength (VSS) and offers supporting information to help analysts assess model fit.

8 **Key words:** initial data analysis; statistical graphics; data visualization; visual inference;
computer vision; machine learning; hypothesis testing; regression analysis;
model diagnostics

9 1. Introduction

10 Regression analysis is a widely used statistical modeling technique for data in many
11 fields. There is a vast array of software for conducting regression modeling and
12 generating diagnostics. The package `lmtest` (Zeileis & Hothorn 2002) provides a suite
13 of conventional tests, while the `stats` package (R Core Team 2022) offers standard
14 diagnostic plots such as residuals vs. fitted values, quantile-quantile (Q-Q) plots,
15 and residuals vs. leverage plots. Additional packages like `jtools` (Long 2022), `olsrr`
16 (Hebbali 2024), `rockchalk` (Johnson 2022), and `ggResidpanel` (Goode & Rey 2019)

¹ Department of Econometrics and Business Statistics, Monash University, Wellington Road, VIC 3800, Australia

² Biological Data Science Institute, The Australian National University, 46 Sullivan's Creek Road, ACT 2600, Australia

³ Department of Statistics, University of Nebraska, Hardin Hall, 3310 Holdrege St Suite 340, Lincoln, NE 68583, United States

Email: patrick.li@anu.edu.au

17 deliver similar graphical diagnostics, often with enhanced aesthetics or interactive
18 features. These tools collectively produce the core diagnostic plots outlined in the
19 classical text by Cook & Weisberg (1982). The `ecostats` package (Warton 2023)
20 extends these diagnostics by incorporating simulation envelopes into residual plots.
21 Meanwhile, DHARMa (Hartig 2022) compares empirical quantiles (0.25, 0.5, and 0.75)
22 of scaled residuals to their theoretical counterparts, with a strong focus on identifying
23 model violations such as heteroscedasticity, misspecified functional forms, and issues
24 specific to generalized linear and mixed-effect models, like over/under-dispersion. It
25 also provides conventional test annotations to reduce the risk of misinterpretation.

26 However, relying solely on subjective assessments of these plots can lead to issues such as
27 over-interpreting random patterns as model violations. Li et al. (2024a) demonstrated
28 that visual inference methods, particularly those using the lineup protocol (Buja
29 et al. 2009), offer more practical and reliable assessments of residual patterns than
30 conventional tests, as they are less sensitive to minor departures. Packages such as
31 `nullabor` (Wickham et al. 2020), `HLMdiag` (Loy & Hofmann 2014), and `regressinator`
32 (Reinhart 2024) support this approach by enabling users to compare observed residual
33 plots with plots generated under null hypothesis, thereby helping to quantify the
34 significance of any detected patterns.

35 As noted in Li et al. (2024b), the lineup protocol has significant limitations in large-
36 scale applications due to its reliance on human labor. To overcome this constraint, a
37 computer vision model was developed alongside a corresponding statistical testing
38 procedure to automate the assessment of residual plots. The model takes as input
39 a residual plot and a set of auxiliary variables (such as the number of observations)
40 and outputs a predicted visual signal strength (VSS). This VSS estimates the degree
41 of deviation between the residual distribution of the fitted model and the reference
42 distribution expected under correct model specification.

43 To make the statistical testing procedure and trained computer vision model widely
44 accessible, we developed the R package `autovi` along with a companion web interface,
45 `autovi.web`, which allows users to automatically assess their residual plots using the
46 trained computer vision model.

47 The remainder of this paper is structured as follows: Section 2 introduces the definition
48 and computation of visual signal strength. Section 4 provides a detailed documentation
49 of the `autovi` package, including its usage and infrastructure. Section 5 focuses on the
50 `autovi.web` interface, describing its design and usage, along with illustrative examples.
51 Finally, Section 6 presents the main conclusions of this work.

52 **2. Definition and computation of visual signal strength**

- 53 To train a computer vision model, a measure of the visible pattern in a plot is needed.
 54 We call this the **visual signal strength** (VSS), which measures how prominently a
 55 specific set of visual patterns appears in an image. This can be computed for a training
 56 set of data, and plots, where the generating distributions are specified.
- 57 In the context of regression model diagnostics, VSS describes the clarity of visual
 58 patterns on a diagnostic plot that may indicate model violations. Violations can be
 59 categorized as weak, moderate, or strong, but here we treat it as a continuous positive
 60 real variable. Importantly, its interpretation depends on how it is linked to a function
 61 of the data or the underlying data generating process. Consequently, the calculation
 62 of VSS can vary across different model classes or within the same model, depending
 63 on the generating function.
- 64 VSS estimates the distance between the residual distribution of a fitted classical normal
 65 linear regression model and a reference distribution (see [Li et al. 2024b](#), for details).
 66 The distance measure is based on the Kullback-Leibler (KL) divergence:

$$D = \log(1 + D_{KL}),$$

67 where D_{KL} is given by:

$$D_{KL} = \int_{\mathbb{R}^n} \log \frac{p(\mathbf{e})}{q(\mathbf{e})} p(\mathbf{e}) d\mathbf{e}, \quad (1)$$

- 68 here, $p(\cdot)$ and $q(\cdot)$ are the probability density functions of the reference residual
 69 distribution P and the true residual distribution Q , respectively.
- 70 This distance measure depends on knowledge of the true residual distribution, which
 71 is unknown in practice. To compute D_{KL} for the training samples, Equation 1 takes
 72 different forms depending on the specific model violations. For instance, where necessary
 73 higher-order predictors, \mathbf{Z} , and their corresponding parameter, $\boldsymbol{\beta}_Z$, are omitted from
 74 the fitted linear model, the distance measure can be expanded as follows:

$$D_{KL} = \frac{1}{2} (\boldsymbol{\mu}_z^\top (\text{diag}(\mathbf{R}\sigma^2))^{-1} \boldsymbol{\mu}_z),$$

75 where $\mu_z = \mathbf{R}\mathbf{Z}\beta_z$, $\mathbf{R} = \mathbf{I}_n - \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top$ and \mathbf{X} is the design matrix of the
 76 regression model.

77 The computer vision model approximates this mapping from a set of residuals to
 78 its corresponding distance measure. It is trained on a large number of synthetic
 79 regression models, each designed to simulate specific violations of classical linear
 80 regression assumptions. These models incorporate non-linearity through Hermite
 81 polynomial transformations of predictors, heteroskedasticity by making the error
 82 variance a predictor-dependent function, and non-normality by drawing residuals from
 83 distributions such as discrete, uniform, and lognormal. Both simple and multiple linear
 84 regression structures are used, with controlled parameters to generate diverse and
 85 complex residual patterns. Since the data-generating process is known, the distance
 86 measure D can be explicitly calculated, enabling supervised training. The computer
 87 vision model takes a residual plot as input and outputs the corresponding distance
 88 measure, learning to quantify model violations directly from visual patterns. Additional
 89 details are provided in [Li et al. \(2024b\)](#).

90 3. Definition and simulation of null and bootstrapped residuals

91 In the subsequent sections, we will frequently refer to null residuals and bootstrapped
 92 residuals, so it is helpful to first define and explain how they are generated.

93 **Null residuals** are used to generate null plots within the lineup protocol framework,
 94 serving as the foundation for the statistical testing in our automated residual plot
 95 assessment. Specifically, they represent residuals generated under the null hypothesis
 96 that the model is correctly specified. A common method for simulating null residuals
 97 in linear regression involves sampling from a normal distribution with mean zero and
 98 variance equal to the estimated variance of the error term. These simulated residuals
 99 and their corresponding plots depict what one would expect from a correctly specified
 100 model. If the true residual plot exhibits noticeable deviations from these null plots, it
 101 may suggest model misspecification.

102 Our computer vision model is trained to assign lower VSS to null plots and higher VSS
 103 to plots that display distinct patterns. Accordingly, statistical testing is performed
 104 by computing the proportion of null plots whose VSS equals or exceeds that of the
 105 observed residual plot. This proportion serves as a p-value for a one-sided hypothesis
 106 test.

107 **Bootstrapped residuals** are obtained by refitting the model on bootstrap samples,
108 which are generated by sampling individual observations with replacement from the
109 original dataset. The residual plots obtained from these refitted models are evaluated
110 using the same computer vision model. The predicted VSS from the bootstrapped
111 plots provide an empirical estimate of the variation in the VSS of the observed
112 residual plot. By examining the proportion of bootstrapped plots that also exhibit
113 significant violations, we can assess whether the original conclusion is robust to
114 sampling variability.

115 4. R package: autovi

116 The main purpose of **autovi** is to provide rejection decisions and p -values for testing
117 the null hypothesis (H_0) that the regression model is correctly specified. The package
118 provides automated interpretation of residual plots using computer vision. The name
119 **autovi** stands for **a**utomated **v**isual **i**nference. This functionality can be accessed
120 through the R package **autovi**, or through a web interface, **autovi.web**, which enables
121 users to perform analyses without installing R, Python, or their associated dependencies
122 locally.

123 4.1. Motivation

124 Figure 1 shows three sets of plots of residuals against fitted values. The simulated
125 example in (a) might be interpreted as a heteroscedastic pattern, however the
126 automated reading would predict this to have a visual signal strength (VSS) of
127 1.53, with a corresponding p -value of 0.25. This means it would be interpreted as
128 a good residual plot, that there is nothing in the data to indicate a violation of
129 model assumptions. Skewness in the predictor variables is generating the apparent
130 heteroscedasticity, where the smaller variance in residuals at larger fitted values is
131 due to smaller sample size only. The Breusch-Pagan test (Breusch & Pagan 1979) for
132 heteroscedasticity would also not reject this as good residual plot.

133 The data in (b) is generated by fitting a linear model predicting `mpg` based on `hp`
134 using the datasets::`mtcars`. It is a small data set, and there is a hint of nonlinear
135 structure not captured by the model. The automated plot reading would predict a
136 VSS of 3.57, which has a p -value less than 0.05. That is, the nonlinear structure is
137 most likely real, and indicates a problem with the model. The conventional test, a
138 Ramsey Regression Equation Specification Error Test (RESET) (Ramsey 1969) would
139 also strongly detect the nonlinearity.

140 The third example is generated using the `surreal` package ([Balamuta 2024](#)), where
 141 structure residuals are embedded in the data. In this case, a quote inspired by Tukey,
 142 “visual summaries focus on unexpected values”, is used to define the residual structure.
 143 The automated plot reading predicts the VSS to be 5.87, with a *p*-value less than
 144 0.05. Visually, the structure is strikingly clear, but a RESET test for nonlinear
 145 structure would not report a problem. (It would be detected by a Breusch-Pagan for
 146 heteroscedasticity and also Shapiro-Wilk test ([Shapiro & Wilk 1965](#)) for non-normality.)

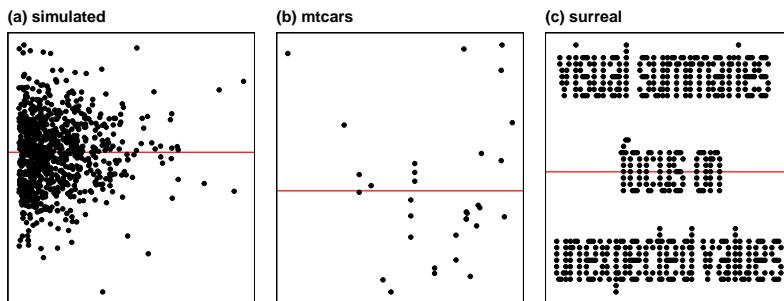


Figure 1. Reading residual plots can be a difficult task, particularly for students new to statistical modeling. The `autovi` package makes it easier. Here are three examples of residual plots, which may appear to have structure. According to `autovi`, the visual signal strengths (VSS) of these three examples are approximately (a) 1.53, (b) 3.57, (c) 5.87, resulting in (b), (c) being significant violations of good residuals, but (a) is consistent with a good residual plot.

147 4.2. Implementation

148 The `autovi` package is built on the `bandicoot` object-oriented programming (OOP)
 149 system ([Li 2024](#)), marking a departure from R’s traditional S3 generic system. This
 150 OOP architecture enhances flexibility and modularity, allowing users to redefine key
 151 functions through method overriding.

152 The `autovi` infrastructure effectively integrates multiple programming languages and
 153 libraries into a comprehensive analytical tool. It relies on five core libraries from
 154 Python and R, each playing a critical role in the analysis pipeline. In Python, `pillow`
 155 ([Clark et al. 2015](#)) handles image processing tasks such as reading and resizing PNG
 156 files of residual plots, then converting them into input tensors for further analysis.
 157 `TensorFlow` ([Abadi et al. 2016](#)), a key component of modern machine learning, is used
 158 to predict the VSS of these plots using a pre-trained convolutional neural network.

159 In the R environment, `autovi` utilizes several libraries. `ggplot2` ([Wickham 2016](#))
 160 generates the initial residual plots, saved as PNG files for visual input. `cassowaryr`

161 (Mason et al. 2022) computes scagnostics (scatter plot diagnostics), providing numerical
 162 features that capture statistical properties of the plots. These scagnostics complement
 163 the visual analysis by offering quantitative metrics as secondary input to the
 164 computer vision model. `reticulate` (Ushey, Allaire & Tang 2024) enables seamless
 165 communication between R and Python.

166 4.3. Installation

167 The `autovi` package is available on CRAN. It is actively developed and maintained,
 168 with the latest updates accessible on GitHub. This paper uses `autovi` version 0.4.2.
 169 The package includes internal functions to check the current Python environment used
 170 by the `reticulate` package. If the necessary Python packages are not installed in the
 171 Python interpreter, an error will be raised. If you want to select a specific Python
 172 environment, you can do so by calling the `reticulate::use_python()` function before
 173 using the `autovi` package.

174 We recommend using the Shiny app `autovi.web` if users encounter installation
 175 problems.

176 4.4. Usage

177 4.4.1. Numerical summary

178 Three steps are needed to get an automated assessment of a set of residuals and fitted
 179 values:

- 180 1. Load the `autovi` package using the `library()` function.
- 181 2. Create a checker object with a linear regression model.
- 182 3. Call the `check()` method of the checker, which, by default, predicts the VSS for
 the true residual plot, 100 null plots, and 100 bootstrapped plots. The method
 stores the predictions internally and prints a concise results report.

185 The code to do this is:

```
library(autovi)
checker <- residual_checker(lm(dist ~ speed, data = cars))
checker$check()
```

186 It produces the following summary:

```
187 <AUTO_VI object>
188 Status:
```

```

189 - Fitted model: lm
190 - Keras model: (None, 32, 32, 3) + (None, 5) -> (None, 1)
191   - Output node index: 1
192 - Result:
193   - Observed visual signal strength: 3.16 (p-value = 0.0396)
194   - Null visual signal strength: [100 draws]
195     - Mean: 1.274
196     - Quantiles:

197
198   25%    50%    75%    80%    90%    95%    99%
199   0.802  1.111  1.575  1.666  1.919  2.657  3.348
200

201 - Bootstrapped visual signal strength: [100 draws]
202   - Mean: 2.795 (p-value = 0.05941)
203   - Quantiles:

204
205   25%    50%    75%    80%    90%    95%    99%
206   2.455  2.941  3.177  3.300  3.474  3.537  3.668
207

208 - Likelihood ratio: 0.7333 (boot) / 0.06284 (null) = 11.67

```

The summary includes observed VSS of the true residual plot and associated p -value of the automated visual test. The p -value is the proportion of null plots (out of the total 100) that have VSS greater than or equal to that of the true residual plot. The report also provides sample quantiles of VSS for null samples and bootstrapped data plots, providing more information about the sampling variability and a likelihood of model violations. The likelihood is computed from the proportion of values greater than the observed VSS in both the bootstrapped data values and the simulated null values.

217 4.4.2. Visual summary

218 Users can visually inspect the original residual plot alongside a sample null plot using
219 `plot_pair()` or a lineup of null plot `plot_lineup()`. This visual comparison can
220 clarify why H_0 is either rejected or not, and help identify potential remedies.

```
checker$plot_pair()
```

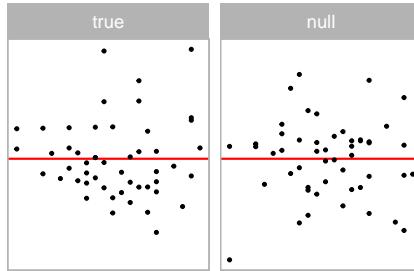


Figure 2. True plot alongside one null plot, for quick comparison.

221 The `plot_pair()` method (Figure 2) displays the true residual plot on the left and a
 222 single null plot on the right. If a full lineup was shown, the true residual plot would
 223 be embedded in a page of null plots. Users should look for any distinct visual patterns
 224 in the true residual plot that are absent in the null plot. Running these functions
 225 multiple times can help any visual suspicions, as each execution generates new random
 226 null plots for comparison.

227 The package offers a straightforward visualization of the assessment result through
 228 the `summary_plot()` function.

```
checker$summary_plot()
```

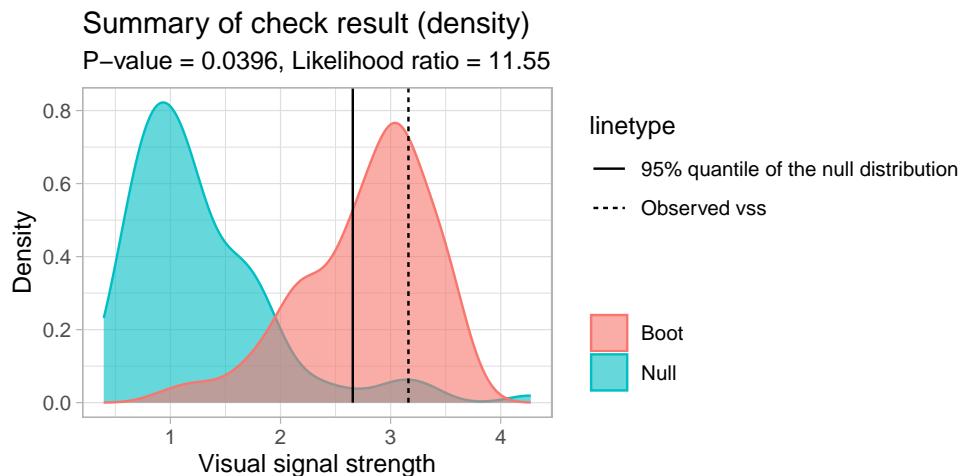


Figure 3. Summary plot comparing the densities of VSS for bootstrapped residual samples (red) relative to VSS for null plots (blue).

- 229 In the result, shown in Figure 3, the blue area represents the density of VSS for null
 230 residual plots, while the red area shows the density for bootstrapped residual plots.
 231 The dashed line indicates the VSS of the true residual plot, and the solid line marks
 232 the critical value at a 95% significance level. The p -value and the likelihood ratio are
 233 displayed in the subtitle. The likelihood ratio represents the ratio of the likelihood
 234 of observing the VSS of the true residual plot from the bootstrapped distribution
 235 compared to the null distribution.
- 236 Interpreting the plot involves several key aspects. If the dashed line falls to the right of
 237 the solid line, it suggests rejecting the null hypothesis. The degree of overlap between
 238 the red and blue areas indicates similarity between the true residual plot and null
 239 plots; greater overlap suggests more similarity. Lastly, the portion of the red area to
 240 the right of the solid line represents the percentage of bootstrapped models considered
 241 to have model violations.
- 242 This visual summary provides an intuitive way to assess the model's fit and potential
 243 violations, allowing users to quickly grasp the results of the automated analysis.

244 4.5. Modularized infrastructure

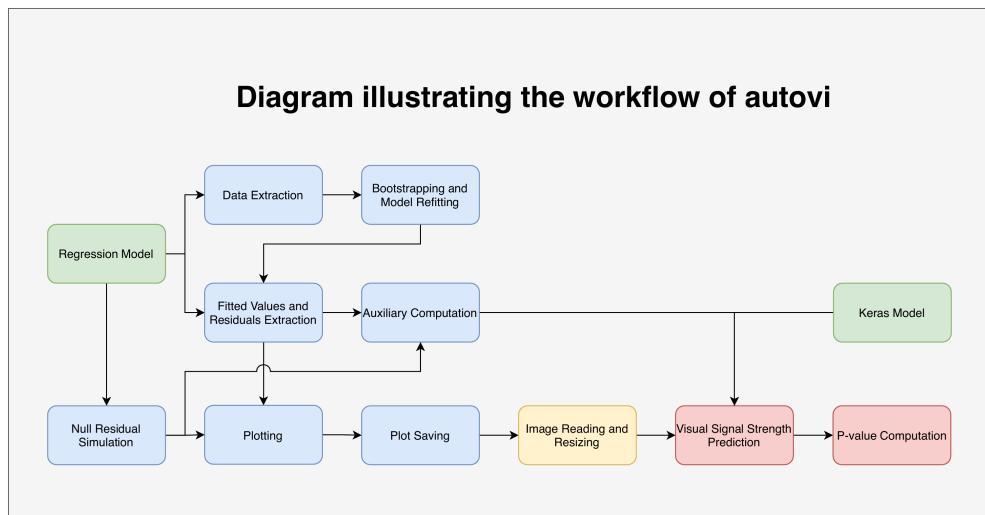


Figure 4. Diagram illustrating the infrastructure of the R package `autovi`. The modules in green are primary inputs provided by users. Modules in blue are overridable methods that can be modified to accommodate users' specific needs. The module in yellow is a pre-defined non-overridable method. The modules in red are primary outputs of the package.

- 245 The initial motivation for developing `autovi` was to create a convenient interface for
 246 sharing the models described and trained in [Li et al. \(2024b\)](#). However, recognizing

247 that the classical normal linear regression model represents a restricted class of
248 models, we sought to avoid limiting the potential for future extensions, whether by
249 the original developers or other developers. As a result, the package was designed to
250 function seamlessly with linear regression models with minimal modification and few
251 required arguments, while also accommodating other classes of models through partial
252 infrastructure substitution. This modular and customizable design allows `autovi` to
253 handle a wide range of residual diagnostics tasks.

254 The infrastructure of `autovi` consists of ten core modules: data extraction,
255 bootstrapping and model refitting, fitted values and residuals extraction, auxiliary
256 computation, null residual simulation, plotting, plot saving, image reading and resizing,
257 VSS prediction, and p -value computation. Each module is designed with minimal
258 dependency on the preceding modules, allowing users to customize parts of the
259 infrastructure without affecting its overall integrity. An overview of this infrastructure
260 is illustrated in Figure 4.

261 The package takes regression models and a `Keras` model as primary inputs. Modules for
262 VSS prediction and p -value computation are fixed but accessible via function arguments,
263 using `TensorFlow` for inference and statistical testing. The image loading module is
264 also fixed, using `PIL` to read and resize images based on the `Keras` model's input shape.
265 The remaining seven modules are overridable, allowing users to adapt the workflow as
266 needed. The data extraction module extracts a `data.frame` containing variables used
267 in the regression model. The bootstrapping and refitting module resamples the data
268 and refits the model. The fitted values and residuals extraction module returns these
269 values as a `data.frame`. The auxiliary computation module calculates scagnostics
270 such as monotonicity. The plotting module generates a `ggplot` in a standard format,
271 and the plot saving module exports it at the same resolution as the training images.
272 These modules are described in detail in the package documentation.

273 4.6. Extension to Other Model Classes

274 The `autovi` R package can be extended to accommodate other classes of models
275 beyond linear regression, such as generalized linear models (`glm`). This is achieved
276 by substituting the relevant overridable modules, and if needed, supplying a different
277 `Keras` model.

278 We provide an example of defining a new checker class tailored for Poisson regression
279 using the `glm` framework:

- 280 1. Define a new class using `new_class()` with `AUTO_VI` as the parent class.
- 281 2. Override the necessary methods using `register_method()`. In this example, we
282 use Pearson residuals. To simulate null residuals, we assume the fitted model
283 is correct and the estimated coefficients are accurate. New response values are
284 generated accordingly, and a new model is fitted to this simulated response. Null
285 residuals are then extracted from this refitted model.
- 286 3. Create an alias for the `instantiate()` method of the new class.

```
AUTO_POIS_VI <- new_class(AUTO_VI, class_name = "AUTO_POIS_VI")
register_method(
  AUTO_POIS_VI,
  get_fitted_and_resid = function(fitted_model = self$fitted_model) {
    tibble(.fitted = fitted(fitted_model),
           .resid = resid(fitted_model, type = "pearson"))
  },
  null_method = function(fitted_model = self$fitted_model) {
    dat <- model.frame(fitted_model)
    dat[[1]] <- rpois(nrow(dat), lambda = fitted(fitted_model))
    new_mod <- update(fitted_model, data = dat)
    return(self$get_fitted_and_resid(new_mod))
  }
)
auto_pois_vi <- AUTO_POIS_VI$instantiate
```

- 287 The resulting checker class can be employed analogously to the linear model case
288 described in Section 4.4. For illustration, we fit a Poisson model in which the quadratic
289 term of the predictor x is intentionally omitted. This misspecification manifests as
290 a pronounced U-shaped pattern in the lineup display (see Figure 5), which is also
291 successfully identified by the computer vision model, yielding a p-value substantially
292 below the conventional threshold of 0.05.

```
x <- rnorm(300, sd = 0.5)
y <- rpois(300, lambda = exp(1 + x + x^2))
pois_checker <- auto_pois_vi(
  glm(y ~ x, family = "poisson"),
  keras_model = get_keras_model("vss_phn_32"))
)
pois_checker$plot_lineup()
```

The true residual plot is at position 4.

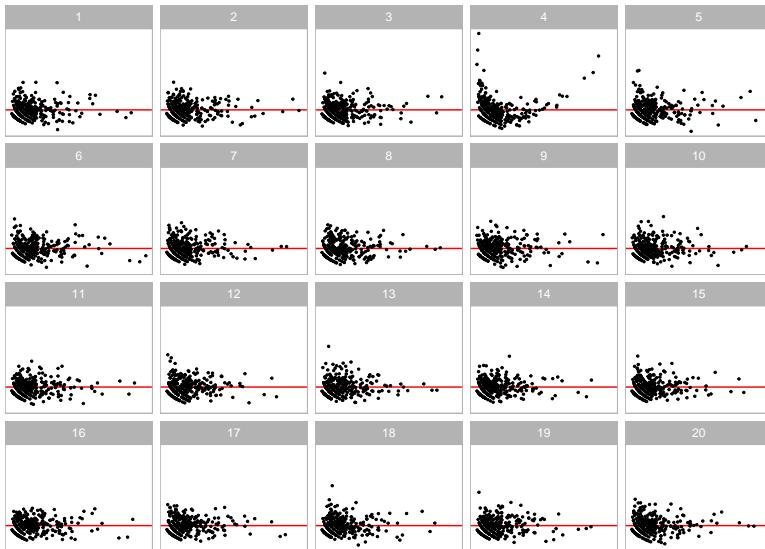


Figure 5. A lineup of residual plots from Poisson generalized linear models, with the true residual plot at position 4, which displays a distinct U-shaped pattern. In contrast, the null plots show characteristics broadly consistent with well-behaved residuals from linear regression models.

```
pois_checker$check()
```

```

293 <AUTO_POIS_VI object>
294 Status:
295   - Fitted model: glm, lm
296   - Keras model: (None, 32, 32, 3) + (None, 5) -> (None, 1)
297     - Output node index: 1
298   - Result:
299     - Observed visual signal strength: 4.875 (p-value = 0.009901)
300     - Null visual signal strength: [100 draws]
301       - Mean: 1.331
302       - Quantiles:
303
304         25%    50%    75%    80%    90%    95%    99%
305         1.035  1.233  1.488  1.644  1.941  2.276  2.639
306
307     - Bootstrapped visual signal strength: [100 draws]
308       - Mean: 5.51 (p-value = 0.009901)

```

309 - Quantiles:

310

25%	50%	75%	80%	90%	95%	99%
5.330	5.505	5.698	5.735	5.830	5.903	6.013

311

312

313

314 - Likelihood ratio: 0.05096 (boot) / 0 (null) = Extremely large

315 It is important to note, however, that the pre-trained computer vision model included
 316 in `autovi`, such as `vss_phn_32` (see `list_keras_models()` for the full list of available
 317 models), was developed specifically for diagnostics of linear regression. Its applicability
 318 to other model classes relies on the assumption that the null residual plots exhibit
 319 characteristics broadly consistent with those of well-behaved linear regression residuals,
 320 that is, residuals should be approximately randomly scattered around zero, display
 321 roughly constant variance across the range of fitted values, and exhibit no discernible
 322 structure or curvature. If these conditions are not met, or if model violations do not
 323 give rise to visually detectable patterns, the validity of the automated diagnostics may
 324 be compromised. In such cases, users are encouraged to train and apply their own
 325 `Keras` models. Detailed guidance on model training and discussion on extending the
 326 methodology to other model classes can be found in [Li et al. \(2024b\)](#).

327 5. Web interface: `autovi.web`

328 The `autovi.web` shiny application extends the functionality of `autovi` by offering a
 329 user-friendly web interface for automated residual plot assessment. This eliminates the
 330 common challenges associated with software installation, so users can avoid managing
 331 Python environments or handling version requirements for R libraries. The platform
 332 is cross-platform and accessible on various devices and operating systems, making it
 333 suitable even for users without R programming experience. Additionally, updates are
 334 managed centrally, ensuring that users always have access to the latest features. This
 335 section discusses the implementation based on `autovi.web` version 0.1.0.

336 5.1. Implementation

337 The interface `autovi.web` is built using the `shiny` ([Chang et al. 2022](#)) and
 338 `shinydashboard` ([Chang & Borges Ribeiro 2021](#)) R packages. Hosted on the
 339 `shinyapps.io` domain, the application is accessible through any modern web browser.
 340 The R packages `htmltools` ([Cheng et al. 2024](#)) and `shinycssloaders` ([Sali & Attali](#)

341 2020) are used to render markdown documentation in shiny application, and for loading
342 animations for shiny widgets, respectively.

343 Determining the best way to implement the backend was difficult. In our initial
344 planning for `autovi.web`, we considered implementing the entire web application using
345 the `webr` framework (Moon 2020), which would have allowed the entire application to
346 run directly in the user’s browser. However, `webr` does not support packages which use
347 compiled fortran code, which is required by `splancs` (Rowlingson & Diggle 2023), a
348 dependency of `autovi`. In the future, it is possible that a working Emscripten (Zakai
349 2011) version of this package may allow full `webr` support.

350 We also explored the possibility of implementing the web interface using frameworks
351 built on other languages, such as Python. However, server hosting domains that
352 natively support Python servers typically do not have the latest version of R installed.
353 Additionally, calling R from Python is typically done using the `rpy2` Python library
354 (Gautier 2024), but this approach can be awkward when dealing with language syntax
355 related to non-standard evaluation. Another option we considered was renting a server
356 where we could have full control, such as those provided by cloud platforms like
357 Google Cloud Platform (GCP) or Amazon Web Services (AWS). However, deploying
358 and maintaining the server securely requires some expertise. Ultimately, the most
359 practical solution was to use the `shiny` and `shinydashboard` frameworks, which are
360 well-established in the R community and offer a solid foundation for web application
361 development.

362 The server-side configuration of `autovi.web` is carefully designed to support its
363 functionality. Most required Python libraries, including `pillow` and `numpy`, are pre-
364 installed on the server. These libraries are integrated into the Shiny application using
365 the `reticulate` package, which provides an interface between R and Python.

366 Due to shinyapps.io’s resource policy, inactive servers enter sleep mode, clearing the
367 local Python environment. When reactivated for a new session, libraries must be
368 reinstalled. While this ensures a clean environment for each session, it may lead to
369 slightly longer loading times for the first user after a period of inactivity.

370 In contrast to `autovi`, `autovi.web` leverages `TensorFlow.js`, a JavaScript library
371 that allows the execution of machine learning models directly in the browser. This
372 choice enables native browser execution, enhancing compatibility across different user
373 environments, and shifts the computational load from the server to the client-side.

374 TensorFlow.js also offers better scalability and performance, especially when dealing
375 with resource-intensive computer vision models on the web.

376 While autovi requires downloading the pre-trained computer vision models from
377 GitHub, these models in “keras” file format are incompatible with TensorFlow.js.
378 Therefore, we extract and store the model weights in JSON files and include
379 them as extra resources in the Shiny application. When the application initializes,
380 TensorFlow.js rebuilds the computer vision model using these pre-stored weights.

381 To allow communication between TensorFlow.js and other components of the Shiny
382 application, the shinyjs R package ([Attali 2021](#)) is used. This package allows calling
383 custom JavaScript code within the Shiny framework. The specialized JavaScript
384 code for initializing TensorFlow.js and calling TensorFlow.js for VSS prediction is
385 deployed alongside the Shiny application as additional resources.

386 5.2. Usage

387 The workflow of autovi.web is designed to be straightforward, with numbered
388 steps displayed in each panel. There are two example datasets provided by the
389 web application. The single residual plot example uses the dino dataset from the
390 R package [datasauRus](#) ([Davies, Locke & D'Agostino McGowan 2022](#)). The lineup
391 example uses residuals from a simulated regression model that has a non-linearity
392 issue. We walk through the lineup example to further demonstrate the workflow of
393 the web application.

394 5.2.1. Reading data and setting parameters

395 The user can select to upload data as either a single set of residuals and fitted values
396 in a two (or more) column CSV file or a pre-computed lineup of residuals and null
397 datasets in a three (or more) column CSV file (i.e. multiple sets of residuals and fitted
398 values with a column indicating the set label). Here we illustrate use with lineup
399 example data sets (Figure 6). To use the lineup example data, click the “Use Lineup
400 Example” button. The data status will then update to show the number of rows and
401 columns in the dataset, and the CSV type will automatically be selected to the correct
402 option. Since the example dataset follows the variable naming conventions assumed
403 by the web application, the columns for fitted values, residuals, and labels of residual
404 plots are automatically mapped such that the column named as `.fitted` is mapped
405 to fitted values, `.resid` is mapped to residuals and if applicable, `.sample` to labels of
406 the residual set (middle image). If the user is working with a custom dataset, these

407 options must be set accordingly. Whenever a data containing a lineup, the user must
 408 manually select the label for the true residual plot, otherwise the web application does
 409 not provide all the results. The last step is to click the play button (right image) to
 410 start the assessment.




Figure 6. To begin the workflow for `autoovi` using the lineup example dataset, the user clicks the “Use Lineup Example” button (left) to load the example dataset, during which the data status and CSV type will be automatically updated. The user must manually select the label for the true residual plot (middle) to compute further results. The user initiates the assessment of the lineup example data by clicking the run button (right).

411 5.2.2. Results provided

412 Results are provided in multiple panels. The first row of the table Figure 7 is the most
 413 crucial to check, as it provides the VSS and the rank of the true residual plot among
 414 the other plots. The summary text beneath the table provides the *p*-value, which can
 415 be used for quick decision-making. The lineup is for manual inspection, and the user
 416 should see if the true residual plot is visually distinguishable from the other plots, to
 417 confirm if the model violation is serious.

418 The density plot in Figure 8 offers a more robust result, allowing the user to compare
 419 the distribution of bootstrapped VSS with the distribution of null VSS. Finally, the
 420 grayscale attention map (right image) can be used to check if the target visual features,
 421 like the non-linearity present in the lineup example, are captured by the computer
 422 vision model, ensuring the quality of the assessment. The attention map is the gradient
 423 of the model output with respect to the grayscale image input, indicating the sensitivity
 424 of the output to each pixel.

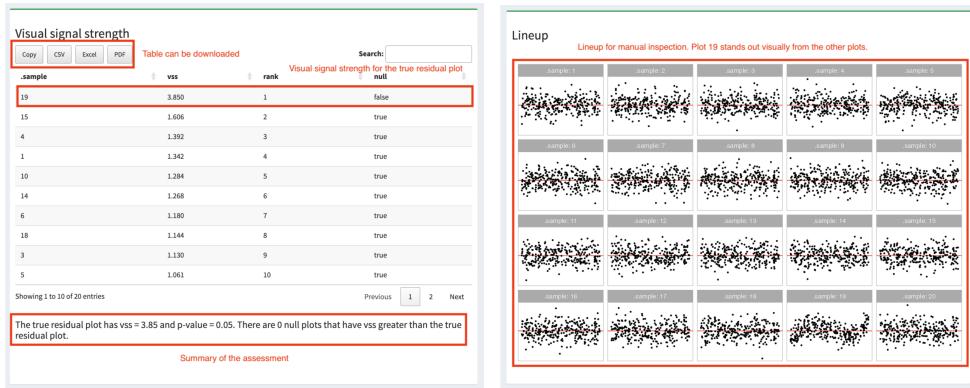


Figure 7. Results for the lineup. The VSS of the true residual plot is displayed in the first row of the table of VSS values for all the null plots (left image), with a summary text beneath the table providing the p -value to aid in decision-making. A lineup of residual plots allows for manual inspection (right image).

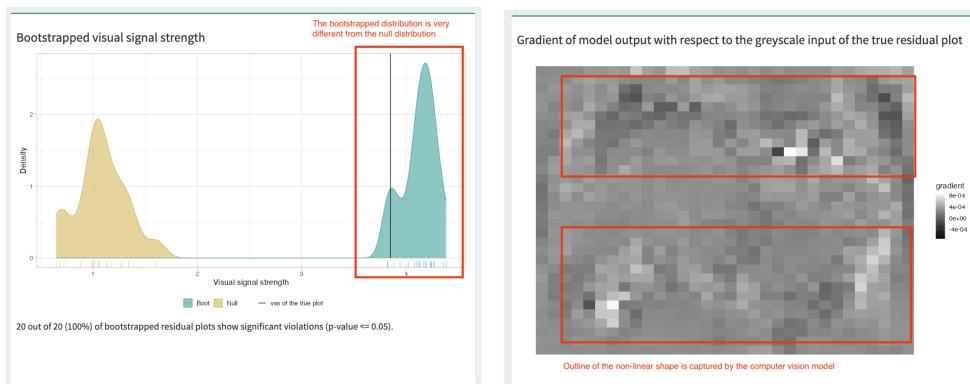


Figure 8. Summaries assessing the strength of the pattern and which elements of the plot contribute. The density plot helps verify if the bootstrapped distribution differs from the null distribution (left image). The attention map (right image) offers insights into whether the computer vision model has captured the intended visual features of the true residual plot.

425

6. Conclusions

426 This paper presents new regression diagnostics software, the R package **autovi** and
 427 its accompanying web interface, **autovi.web**. It addresses a critical gap in the current
 428 landscape of statistical software. While regression tools are widely available, effective
 429 and efficient diagnostic methods have lagged behind, particularly in the field of residual
 430 plot interpretation.

431 The **autovi** R package, introduced in this paper, automates the assessment of residual
 432 plots by incorporating a computer vision model, reducing reliance on time-consuming

433 and potentially inconsistent human interpretation. This automation improves the
434 efficiency of the diagnostic process and promotes consistency in model evaluation
435 across different users and studies.

436 The development of the accompanying Shiny app, `autovi.web`, expands access to these
437 advanced diagnostic tools, by providing a user-friendly interface. It makes automated
438 residual plot assessment accessible to a broader audience, including those who may not
439 have extensive programming experience. This web-based solution effectively addresses
440 the potential barriers to adoption, such as complex dependencies and installation
441 requirements, that are often associated with advanced statistical software.

442 The combination of `autovi` and `autovi.web` offers a comprehensive solution to the
443 challenges of residual plot interpretation in regression analysis. These tools have the
444 potential to significantly improve the quality and consistency of model diagnostics
445 across various fields, from academic research to industry applications. By automating
446 a critical aspect of model evaluation, they allow researchers and analysts to focus more
447 on interpreting results and refining models, rather than grappling with the intricacies
448 of plot assessment.

449 The framework established by `autovi` and `autovi.web` opens up exciting possibilities
450 for further research and development. Future work could explore the extension of these
451 automated assessment techniques to other types of diagnostic plots and statistical
452 models, potentially revolutionizing how we approach statistical inference using visual
453 displays more broadly.

454 7. Resources and supplementary material

455 The current version of `autovi` can be installed from CRAN, and source
456 code for both packages are available at github.com/TengMCing/autovi and
457 github.com/TengMCing/autovi_web respectively. The web interface is available from
458 autoviview.netlify.app.

459 This paper is reproducibly written using Quarto ([Allaire et al. 2024](#)) powered by
460 Pandoc ([MacFarlane, Krewinkel & Rosenthal 2024](#)) and pdfTeX. The full source code
461 to reproduce this paper is available at github.com/TengMCing/autovi_paper.

462 These R packages were used for the work: `tidyverse` ([Wickham et al. 2019](#)), `lmtest`
463 ([Zeileis & Hothorn 2002](#)), `kableExtra` ([Zhu 2021](#)), `patchwork` ([Pedersen 2022](#)),
464 `rcartocolor` ([Nowosad 2018](#)), `glue` ([Hester & Bryan 2022](#)), `here` ([Müller 2020](#)),

465 `magick` (Ooms 2023), `yardstick` (Kuhn, Vaughan & Hvitfeldt 2024) and `reticulate`
 466 (Ushey, Allaire & Tang 2024).

467 **References**

- 468 ABADI, M., AGARWAL, A., BARHAM, P., BREVDO, E., CHEN, Z., CITRO, C., CORRADO, G.S.,
 469 DAVIS, A., DEAN, J., DEVIN, M. et al. (2016). Tensorflow: Large-scale machine learning on
 470 heterogeneous distributed systems. *arXiv preprint arXiv:1603.04467* .
- 471 ALLAIRE, J., TEAGUE, C., SCHEIDECKER, C., XIE, Y. & DERVIEUX, C. (2024). Quarto. doi:
 472 10.5281/zenodo.5960048. URL <https://github.com/quarto-dev/quarto-cli>.
- 473 ATTALI, D. (2021). *shinyjs: Easily Improve the User Experience of Your Shiny Apps in Seconds*.
 474 URL <https://CRAN.R-project.org/package=shinyjs>. R package version 2.1.0.
- 475 BALAMUTA, J.J. (2024). *surreal: Create Datasets with Hidden Images in Residual Plots*. URL
 476 <https://CRAN.R-project.org/package=surreal>. R package version 0.0.1.
- 477 BREUSCH, T.S. & PAGAN, A.R. (1979). A simple test for heteroscedasticity and random coefficient
 478 variation. *Econometrica: Journal of the Econometric Society* , 1287–1294.
- 479 BUJA, A., COOK, D., HOFMANN, H., LAWRENCE, M., LEE, E.K., SWAYNE, D.F. & WICKHAM, H.
 480 (2009). Statistical inference for exploratory data analysis and model diagnostics. *Philosophical
 481 Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* **367**,
 482 4361–4383.
- 483 CHANG, W. & BORGES RIBEIRO, B. (2021). *shinydashboard: Create Dashboards with 'Shiny'*.
 484 URL <https://CRAN.R-project.org/package=shinydashboard>. R package version 0.7.2.
- 485 CHANG, W., CHENG, J., ALLAIRE, J., SIEVERT, C., SCHLOERKE, B., XIE, Y., ALLEN, J.,
 486 MCPHERSON, J., DIPERT, A. & BORGES, B. (2022). *shiny: Web Application Framework for
 487 R*. URL <https://CRAN.R-project.org/package=shiny>. R package version 1.7.3.
- 488 CHENG, J., SIEVERT, C., SCHLOERKE, B., CHANG, W., XIE, Y. & ALLEN, J. (2024). *htmltools:
 489 Tools for HTML*. URL <https://CRAN.R-project.org/package=htmltools>. R package version
 490 0.5.8.
- 491 CLARK, A. et al. (2015). Pillow (pil fork) documentation. *readthedocs* .
- 492 COOK, R.D. & WEISBERG, S. (1982). *Residuals and influence in regression*. New York: Chapman
 493 and Hall.
- 494 DAVIES, R., LOCKE, S. & D'AGOSTINO McGOWAN, L. (2022). *datasauRus: Datasets from the
 495 Datasaurus Dozen*. URL <https://CRAN.R-project.org/package=datasauRus>. R package
 496 version 0.1.6.
- 497 GAUTIER, L. (2024). *Python interface to the R language (embedded R)*. URL [https://pypi.org/project/rpy2/](https://pypi.org/

 498 project/rpy2/). Version 3.5.16.
- 499 GOODE, K. & REY, K. (2019). *ggResidpanel: Panels and Interactive Versions of Diagnostic Plots
 500 using 'ggplot2'*. URL <https://CRAN.R-project.org/package=ggResidpanel>. R package version
 501 0.3.0.
- 502 HARTIG, F. (2022). *DHARMa: Residual Diagnostics for Hierarchical (Multi-Level / Mixed)
 503 Regression Models*. URL <https://CRAN.R-project.org/package=DHARMa>. R package
 504 version 0.4.6.
- 505 HEBBALI, A. (2024). *olsrr: Tools for Building OLS Regression Models*. URL [https://CRAN.R-project.org/package=olsrr](https://CRAN.R-

 506 project.org/package=olsrr). R package version 0.6.0.
- 507 HESTER, J. & BRYAN, J. (2022). *glue: Interpreted String Literals*. URL [https://CRAN.R-project.org/package=glue](https://CRAN.R-

 508 project.org/package=glue). R package version 1.6.2.
- 509 JOHNSON, P.E. (2022). *rockchalk: Regression Estimation and Presentation*. URL [https://CRAN.R-project.org/package=rockchalk](https://CRAN.R-

 510 project.org/package=rockchalk). R package version 1.8.157.

- 511 KUHN, M., VAUGHAN, D. & HVITFELDT, E. (2024). *yardstick: Tidy Characterizations of Model*
 512 *Performance*. URL <https://CRAN.R-project.org/package=yardstick>. R package version 1.3.1.
- 513 LI, W. (2024). *bandicoot: Light-weight python-like object-oriented system*. URL <https://CRAN.R->
 514 [project.org/package=bandicoot](https://CRAN.R-project.org/package=bandicoot).
- 515 LI, W., COOK, D., TANAKA, E. & VANDERPLAS, S. (2024a). A plot is worth a thousand tests:
 516 Assessing residual diagnostics with the lineup protocol. *Journal of Computational and*
 517 *Graphical Statistics* **33**, 1497–1511. doi:10.1080/10618600.2024.2344612.
- 518 LI, W., COOK, D., TANAKA, E., VANDERPLAS, S. & ACKERMANN, K. (2024b). Automated
 519 assessment of residual plots with computer vision models. *arXiv preprint arXiv:2411.01001* .
- 520 LONG, J.A. (2022). *jtools: Analysis and Presentation of Social Scientific Data*. URL <https://cran.r->
 521 [project.org/package=jtools](https://cran.r-project.org/package=jtools). R package version 2.2.0.
- 522 LOY, A. & HOFMANN, H. (2014). *Hlmdiag: A suite of diagnostics for hierarchical linear models in*
 523 *r*. *Journal of Statistical Software* **56**, 1–28.
- 524 MACFARLANE, J., KREWINKEL, A. & ROSENTHAL, J. (2024). *Pandoc*. URL <https://github.com/>
 525 [jgm/pandoc](https://github.com/jgm/pandoc).
- 526 MASON, H., LEE, S., LAA, U. & COOK, D. (2022). *cassowaryr: Compute Scagnostics on Pairs of*
 527 *Numeric Variables in a Data Set*. URL <https://CRAN.R-project.org/package=cassowary>. R
 528 package version 2.0.0.
- 529 MOON, K.W. (2020). *webr: Data and Functions for Web-Based Analysis*. URL <https://CRAN.R->
 530 [project.org/package=webr](https://CRAN.R-project.org/package=webr). R package version 0.1.5.
- 531 MÜLLER, K. (2020). *here: A simpler way to find your files*. URL <https://CRAN.R-project.org/>
 532 [package=here](https://CRAN.R-project.org/package=here). R package version 1.0.1.
- 533 NOWOSAD, J. (2018). 'CARTOCOLORs' palettes. URL <https://nowosad.github.io/rkartocolor>. R
 534 package version 1.0.
- 535 OOMS, J. (2023). *magick: Advanced Graphics and Image-Processing in R*. URL <https://CRAN.R->
 536 [project.org/package=magick](https://CRAN.R-project.org/package=magick). R package version 2.7.4.
- 537 PEDERSEN, T.L. (2022). *patchwork: The composer of plots*. URL <https://CRAN.R-project.org/>
 538 [package=patchwork](https://CRAN.R-project.org/package=patchwork). R package version 1.1.2.
- 539 R CORE TEAM (2022). *R: A Language and Environment for Statistical Computing*. R Foundation
 540 for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- 541 RAMSEY, J.B. (1969). Tests for specification errors in classical linear least-squares regression
 542 analysis. *Journal of the Royal Statistical Society: Series B (Methodological)* **31**, 350–371.
- 543 REINHART, A. (2024). *regressinator: Simulate and Diagnose (Generalized) Linear Models*. URL
 544 <https://CRAN.R-project.org/package=regressinator>. R package version 0.2.0.
- 545 ROWLINGSON, B. & DIGGLE, P. (2023). *splancs: Spatial and Space-Time Point Pattern Analysis*.
 546 URL <https://CRAN.R-project.org/package=splancs>. R package version 2.01-44.
- 547 SALI, A. & ATTALI, D. (2020). *shinycssloaders: Add Loading Animations to a 'shiny' Output*
 548 *While It's Recalculating*. URL <https://CRAN.R-project.org/package=shinycssloaders>. R
 549 package version 1.0.0.
- 550 SHAPIRO, S.S. & WILK, M.B. (1965). An analysis of variance test for normality (complete samples).
 551 *Biometrika* **52**, 591–611.
- 552 USHEY, K., ALLAIRE, J. & TANG, Y. (2024). *reticulate: Interface to 'Python'*. URL <https://CRAN.R->
 553 [project.org/package=reticulate](https://CRAN.R-project.org/package=reticulate). R package version 1.35.0.
- 554 WARTON, D.I. (2023). Global simulation envelopes for diagnostic plots in regression models. *The*
 555 *American Statistician* **77**, 425–431.
- 556 WICKHAM, H. (2016). *ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York.
 557 URL <https://ggplot2.tidyverse.org>.
- 558 WICKHAM, H., AVERICK, M., BRYAN, J., CHANG, W., McGOWAN, L.D., FRANÇOIS, R.,
 559 GROLEMUND, G., HAYES, A., HENRY, L., HESTER, J., KUHN, M., PEDERSEN, T.L., MILLER,

- 560 E., BACHE, S.M., MÜLLER, K., OOMS, J., ROBINSON, D., SEIDEL, D.P., SPINU, V.,
561 TAKAHASHI, K., VAUGHAN, D., WILKE, C., WOO, K. & YUTANI, H. (2019). Welcome to
562 the tidyverse. *Journal of Open Source Software* **4**, 1686. doi:10.21105/joss.01686.
- 563 WICKHAM, H., CHOWDHURY, N.R., COOK, D. & HOFMANN, H. (2020). *nullabor: Tools for*
564 *Graphical Inference*. URL <https://CRAN.R-project.org/package=nullabor>. R package version
565 0.3.9.
- 566 ZAKAI, A. (2011). Emscripten: an llvm-to-javascript compiler. In *Proceedings of the ACM*
567 *international conference companion on Object oriented programming systems languages and*
568 *applications companion*. pp. 301–312.
- 569 ZEILEIS, A. & HOTHORN, T. (2002). Diagnostic checking in regression relationships. *R News* **2**,
570 7–10.
- 571 ZHU, H. (2021). *kableExtra: Construct complex table with kable and pipe syntax*. URL
572 <https://CRAN.R-project.org/package=kableExtra>. R package version 1.3.4.